Certifying Digital Rights' Expression Languages

by

Bahman Sistany

Ph.D.Thesis Proposal For the Ph.D. degree in
Computer Science

School of Electrical Engineering and Computer Science
Faculty of Engineering
University of Ottawa

Ph.D.Thesis Advisor: Amy Felty

Abstract

Acknowledgements

Dedication

Table of Contents

List of Tables

List of Figures

Chapter 1

Introduction

1.1 Digital Rights

Digital Rights Management (DRM) refers to the digital management of rights associated with the access or usage of digital assets. There are various aspects of rights management however. According to the authors of the whitepaper "A digital rights management ecosystem model for the education community,"([3]) digital rights management systems cover the following four areas: 1) defining rights 2) distributing/acquiring rights 3) enforcing rights and 4) tracking usage.

1.2 Policy Expression Languages

Rights Expression Languages (REL)s, or more precisely Digital Rights Expression Languages (DREL)s when dealing with digital assets deal with the "rights definition" aspect of the DRM ecosystem. A DREL, allows the expression and definition of digital asset usage rights such that other areas of the DRM ecosystem, namely the enforcement mechanism and the usage tracking components can function correctly.

Currently the most popular RELs are the eXtensible rights Markup Language (XrML) [10], and the Open Digital Rights Language (ODRL) [6]. Both of these languages are XML based and are considered declarative languages. XrML has been selected to be the REL for *MPEG-21* which is an ISO standard for multimedia applications. ODRL is also a standards based REL which has been accepted as part of the W3C community with the mandate of standardizing how rights and policies, related to the usage of digital content on the Open Web Platform, *OWP*, are expressed [REF][wikipedia]. ODRL 2.0 supports expression of rights and also privacy rules for social media while ODRL 1.0 was only dealing with the mobile ecosystem – ODRL 1.0 was adopted by the Open Mobile Alliance, *OMA* in 2000.

As popular as both XrML and ODRL are, their adoption and usage is still somewhat limited in practice. Both Apple and Microsoft for example have defined their own

lightweight RELs [REF][problem with RELs paper] in *Fair Play* (Apple) and in *PlayReady* (Microsoft). The authors of [REF][the problem with RELs] argue that both these RELs and other ones are simply too complex to be used effectively since they try to cover much of the DRM ecosystem.

Rights expressions in DRELs and specifically ODRL are used to arbitrate access to assets under conditions. The main construct in ODRL is the *agreement* which specifies users, asset(s) and policies whereby controls on users' access to the assets are described. This is very similar to how access control conditions are expressed in access control policy languages such as eXtensible Access Control Markup Language (XACML) [REF] and SELinux [REF] and therefore we will be generalizing the concept of a policy language throughout this thesis.

1.3 Semantics Of Policies

Formal methods help ensure that a system behaves correctly with respect to a specification of its desired behavior [REF][TAPL]. This specification of the desired behavior is what's referred to as *semantics* of the system. Using formal methods requires defining precise and formal semantics, without which analysis and reasoning about properties of the system in question would become impossible. For example, an issue with the current batch of RELs are due to their semantics being expressed in a natural language (e.g. English) which by necessity results in ambiguous and open to interpretation behavior.

To formalize the semantics of policy languages several approaches have been attempted by various authors. Most are logic based [REF] while others are based on model-checking [REF], operational semantics based interpreters [REF] and web ontology (from the Knowledge Representation Field)[REF]. In this thesis we will focus on the logic based approach to formalizing semantics and will study a specific logic based language that is a translation from a subset of ODRL.

1.4 Logic Based Semantics

Formal logic can represent the statements and facts we express in a natural language like English. Propositional logic is expressive enough to express simple facts as propositions and uses connectives to allow for the negation, conjunction and disjunction of the facts. However propositional logic is not expressive enough to express policies of the kind used in languages like ODRL and XrML. For example, a simple policy expressed in English like "All who pay 5 dollars can watch the movie Toy Story" cannot be expressed in propositional logic because the concept of variables doesn't exist in propositional logic.

A richer logic such as "Predicate Logic" or "First Order Logic" (*FOL*) is more suitable and has the expressive power to represent policies written in English. Moreover, FOL can be used to capture the meaning of policies in an unambiguous way.

2

Halpern and Weissman [REF][Using First Order Logic to Reason about Policies] propose a fragment of FOL to represent and reason about policies. The fragment of FOL they arrive at is called *Lithium* which is decidable and allows for efficiently answering interesting queries. Lithium restricts policies to be written based on the concept of "bipolarity" which disallows by construction policies that both permit and deny an action on an object. Pucella and Weissman [8] specify a predicate logic based language that represents a subset of ODRL.

1.5   Specific Problem

Policy languages and the agreements written in those languages are meant to implement specific goals such as limiting access to specific assets. The tension in designing a policy language is usually between how to make the language expressive enough, such that the design goals for the policy language may be expressed, and how to make the policies verifiable with respect to the stated goals.

As stated earlier, an important part of fulfilling the verifiability goal is to have formal semantics defined for policy languages. For ODRL, authors of [8] have defined a formal semantics based on which they declare and prove a number of important theorems (their main focus is on stating and proving algorithm complexity results). However as with many paper-proofs, the language used to do the proofs while mathematical in nature, uses a lot of intuitive justifications to show the proofs. As such these proofs are difficult to verify or more importantly to "derive". Furthermore the proofs can not be used directly to render a decision on a sample policy (e.g. whether to allow or deny access to an asset). Of course one may (carefully) construct a program based on these proofs for practical purposes but we will have no way of certifying those programs correct, even assuming the original proofs were in fact correct.

While there are paper-proofs for ODRL, as far as we know, similar paper-proofs do not exist for an important (mandatory) access-control policy system, namely SELinux[REF]. In particular no formal proofs (paper based or otherwise) of decidablity of SELinux policies exist in the literature.

1.6   Contributions

In this thesis we will build a language representation framework based on ODRL and definitions in [8]. The framework will be in *Coq* [4] which is both a programming language and a proof-assistant. We will declare and prove decidablity results of subsets of ODRL all the way up to the complete ODRL fragement defined in [8]. We will extract programs from the proofs and demonstrate how they can be used on specific policies to render a specific decision such as "a conflict has been detected".

Beside "certified decidablity results" for ODRL, we will investigate decidablity for SELinux policies, proving decidablity or show why a proof is not possible (if that is the case) and provide proposals to make the policy language decidable.

By using the Coq framework originally built for ODRL to encode and verify agreements written in a second policy language (albeit a different class of policy language: REL vs access-control) we will demonstrate the suitability of this Coq based framework for other policy languages such as XACML[REF].

1.7 What specific work has been accomplished until this point in time? what results were obtained so far?

The encodings for a subset of ODRL which we call ODRL0 (see 2.2) plus some important functions implementing some of the algorithms in [8] have been implemented in Coq. Some of the intermediate theorems have been also been defined and proved.

1.8 What remains to be done to complete the thesis research?

The main decidablity result and its proof for ODRL0 will be completed first. We will add the remaining ODRL constructs incrementally while maintaining decidablity for the main decision algorithm. The remaining constructs include a trouble-some construct ([8]), namely *not[policySet]*. We will show this construct does not change the decidablity result already established.

ODRL0 is enough to be used as a basis for SELinux policies without *constrain*s. SELinux constrains are extra conditions that need to be satisfied (in addition to policies) in order for a permission to be granted. We will investigate decidablity for this subset first. We will then add constrains to the ODRL0 subset (as pre-requisites) and investigate decidablity.

1.9 What is the timetable to complete the work?

The study plan calls for August of 2015 for everything to be completed.

Chapter 2

ODRL0 Syntax

2.1 Introduction

Authors of [8] use abstract syntax instead of XML (ODRL 2.0 can also be encoded in *JSON* and *RDF/OWL Ontology*) to express statements in the ODRL language. Abstract syntax is a more compact representation than XML in which ODRL policies may be written in and furthermore abstract syntax simplifies specifying the semantics as we shall see later. As an example the agreement "If Mary Smith pays five dollars, then she is allowed to print the eBook 'Treasure Island' twice and she is allowed to display it on her computer as many times as she likes" written in ODRL's XML encoding and the equivalent one expressed in the abstract syntax [8].

Listing 2.1: agreement for Mary Smith in XML

```
<agreement>
 <asset> <context> <uid> Treasure Island </uid> </context> </
    asset>
 <permission>
   <display>
    <constraint>
     <cpu> <context> <uid> Mary's computer </uid> </context> <
        /cpu>
    </constraint>
   </display>
   <print>
    <constraint> <count> 2 </count> </constraint>
   </print>
  <requirement>
   <prepay>
    <payment> <amount currency="AUD"> 5.00</amount> </payment>
   </prepay>
  </requirement>
```

```
  </permission>
 <party> <context> <name> Mary Smith </name> </context> </
    party>
</agreement>
```

The agreement in listing 2.1 is shown below using the syntax from [8].

Listing 2.2: agreement for Mary Smith as BNF (as used in [8])

```
agreement
 for Mary Smith
 about Treasure Island
 with prePay[5.00] -> and[cpu[Mary's Computer] => display,
                                    count[2] => print].
```

In the following we will cover the *abstract syntax* of a subset of ODRL that we later
express using Coq's constructs such as *Inductive Types* and Definitions. We will call this
subset *ODRL0* both because it is a variation of Pucella's ODRL language and also because
it is missing some constructs from Pucella's ODRL (and hence from ODRL proper).

2.2   ODRL0

In ODRL0, agreements and facts (i.e. environments) will only contain the number of times
each policy has been used to justify an action. In ODRL0 agreements and facts will not
contain:

1. Which payments have been made

2. Which acknowledgments have been made

This means *Paid* and *Attributed* predicates are not used in ODRL0. Also removed are
related constructs *prepay* and *attribution*. We also had to remove two other constructs
based on *prepay* and *attribution* out of ODRL0 in *inSeq* and *anySeq*. *prepay*, *attribution*,
inSeq and *anySeq* make up what is called *requirements* in ODRL.

In ODRL a *prerequisite* is either *true*, a *constraint*, a *requirement* or a *condition*. *true*
is the prerequisite that always holds. Constraints are facts that are outside of control of
users. For example, there is nothing *Alice* can do to satisfy the constraint "user must be
Bob". Requirements are facts that are in users' control. For example, *Alice* may satisfy the
requirement "The user must pay 5 dollars". Finally conditions are constraints that must
not hold.

In ODRL0, a prerequisite is either *true*, a *constraint*, or the negation of a *constraint*.
So we have removed requirements from the picture and don't have explicit conditions.
Conditions are replaced by a category called *NotCons* directly in the production for pre-
requisites (see listing 2.12). Note that we have also removed the condition *not[policySet]*

6

from ODRL since the authors in [8] have shown the semantics of this component are not well-defined and including it leads to intractability results.

We will add the missing pieces as described above (except for *not[policySet]*) making up what we will call *ODRL1* and perhaps *ODRL2* (the latter only if needed). We will also describe ODRL0 in a *BNF* grammar that looks more like Pucella's ODRL grammar. BNF style grammars are less formal as they give some suggestions about the surface syntax of expressions [Pierce1] without getting into lexical analysis and parsing related aspects such as precedence order of operators. The Coq version in contrast is more formal and could be directly used for building compilers and interpreters. We will present both the BNF version and the Coq version for each construct of ODRL0 [Pierce1]. To get started let's see what the listing 2.2 would look like in ODRL0's Coq version.

Listing 2.3: Coq version of agreement for Mary Smith

```
Agreement (Single MarySmith) Treasure Island
 (PrimitivePolicySet (Constraint (PrePay 5.00))
  (AndPolicy
   (NewList (PrimitivePolicy (Constraint
                      (Principal
                       (Single MarysComputer))) id1 Display)
   (Single (PrimitivePolicy (Constraint (Count 2)) id2 Print))))).
```

2.3   Productions

The top level ODRL0 production is the *agreement*. An agreement expresses what actions a set of subjects may perform on an object and under what conditions. Syntactically an agreement is composed of a set of subjects/users called a *principal* or *prin*, an *asset* and a *policySet*.

Listing 2.4: agreement

```
<agreement> ::= 'agreement' 'for' <prin> 'about' <asset> 'with' <policySet
   >
```

Principals or prins are composed of *subjects* which are specified based on the application e.g. Alice, Bob, etc for the application we will be using throughout.

Listing 2.5: prin

```
<prin> ::=  { <subject₁>, ..., <subjectₘ> }
```

Listing 2.6: subject

```
<subject> ::= N
```

Assets are also application specific but similar to subjects we will use specific ones for the DRM application (taken from [8]). *ebook*, *The Report* and *latestJingle* are examples of specific subjects we will be using throughout. Syntactically an asset is represented as a natural number (N). Similarly for subjects.

Listing 2.7: asset

```
<asset> ::= N
```

Agreements include policy sets. Each policy set specifies a *prerequisite* and a *policy*. In general if the prerequisite holds the policy is taken into consideration. Otherwise the policy will not be looked at. Some policy sets are specified as *exclusive*. The *Primitive Exclusive Policy Sets* are exclusive to agreement's users in that only those users may perform the actions specified in the policy set. The implication is that all other users who are not specified in the agreement's principal (prin) are forbidden from performing the specified actions. Finally policy sets could be grouped together in a *conjunction* allowing a single agreement to be associated with many policy sets.

Listing 2.8: policySet

```
<policySet> ::=
   <preRequisite> → <policy>              ; primitive policy set
   <preRequisite> ↦ <policy>              ; primitive exclusive policy set
   'and'[ <policySet₁>, ..., <policySetₘ> ]
   ; conjunction
```

A policy specifies an action to be performed on an asset, depending of whether the policy's prerequisite holds or not. If the prerequisite holds the agreement's user is permitted to perform the action on the agreement's asset; otherwise permission is denied. Similar to policy sets, policies could also be grouped together in a conjunction. The policy also includes a unique identifier. The policy identifier is added to help the translation (from agreements to formulas) but is optional in ODRL proper.

Listing 2.9: policy

```
<policy> ::=
   <preRequisite> ⇒<policyId> <act>       ; primitive policy
   'and'[ <policy₁>, ..., <policyₘ> ]     ; conjunction
```

An *Action* (*act*) is represented as a natural number. Similar to assets and subjects, actions are application specific. Some example actions taken from [8] are *Display* and *Print*.

Listing 2.10: act

```
<act> ::= N
```

A *Policy Id* (*policyId*) is a unique identifier specified as (increasing) positive integers.

Listing 2.11: policyId

```
<policyId> ::= N
```

In ODRL0 a *prerequisite* is either true or it is a *constraint*. The *true* prerequisite always holds. A constraint is an intrinsic part of a policy and cannot be influenced by agreement's user. Minimum height requirements for popular attractions and rides are examples of we would consider a constraint. The constraint *ForEachMember* is interesting in its expressive power but has complicated semantics as we shall see in the **??** section. Roughly speaking, ForEachMember takes a prin (a list of subjects) and a list L of constraints. The ForEachConstraint holds if each subject in prin satisfies each constraint in L.*NotCons* is a negation of a constraint. The set of prerequisites are closed under conjunction (*AndPrqs*), disjunction (*OrPrqs*) and exclusive disjunction (*XorPrqs*).

Listing 2.12: preRequisite

```
<preRequisite> ::=
   'True'                                    ; always true
   <constraint>                              ; constraint
   'ForEachMember' [<prin> ; <constraint_1>, ..., <constraint_m> ]
   ; constraint distribution
   'not' [ <constraint> ]                    ; suspending constraint
   'and'[ <preRequisite_1>, ..., <preRequisite_m> ]
   ; conjunction
   'or'[ <preRequisite_1>, ..., <preRequisite_m> ]
   ; disjunction
   'xor'[ <preRequisite_1>, ..., <preRequisite_m> ]
   ; exclusive disjunction
```

Constraints are either *Principal, Count* or *CountByPrin*. Principal constraints basically require matching to specified prins. For example, the user being Alice is a Principal constraint. A count constraint refers to a set of policies P and specifies the number of times the user of an agreement has invoked the policies in P to justify her actions. If the count constraint is part of a policy then the set P is composed of the single policy. In the case that the count constraint is part of a policy set, the set P is the set of policies specified in the policy set.

Listing 2.13: constraint

```
<constraint> ::=
   <prin>                                    ; principal
   'Count' [N]                               ; number of executions
   <prin> ('Count' [N])                      ; number of executions by prin
```

Chapter 3

ODRL0 Syntax In Coq

3.1 Introduction to Coq

Coq is known first and foremost as a proof-assistant. The underlying formal language that Coq uses is a much more expressive version of typed lambda calculus called *Calculus of (Co)Inductive Constructions* or *CIC* where proofs and programs can both be represented. For example, CIC adds polymorphisn (terms depending on types), type operators (types depending on types) and dependent types (types depending on terms).

Specifications of programs in Coq may be expressed using the specification language *Gallina* [5]. Coq is then used to develop proofs to show that a program's run-time behavior satisfies its specification. Such programs are called *certified* because they are formally verified and confirmed to conform to their specifications [2].

Assertions or propositions are statements about values in Coq such as $3 < 8$ or $8 < 3$ that may be true, false or even be only conjectures. To verify that a proposition is true a proof needs to constructed. While paper-proofs use a combination of mathematics and natural language to describe their proofs, Coq provides a formal (and therefore unambiguous) language that is based on proof-theory to develop proofs in. Verification of complex proofs is possible because one can verify the intermediate proofs or sub-goals in steps, each step being derived from the previous by following precise derivation rules. The Coq proof engine solves successive goals by using predefined *tactics*. Coq tactics are commands to manipulate the local context and to decompose a goal into simpler goals or sub-goals [2].

3.2 ODRL0 Syntax

ODRL0 productions were presented as high level abstract syntax in 2.3. Below we present the corresponding encodings in Coq.

An agreement is a new inductive type in Coq by the same name. The constructor *Agreement* takes a *prin*, an *asset* and a *policySet*. *prin* is defined to be a non empty list of *subject*s (see listing 3.1).

Types *asset*, *subject*, *act* and *policyId* are simply defined as *nat* which is the datatype of natural numbers defined in coq's library module *Coq.Init.Datatypes* (*nat* is itself an inductive datatype). We use Coq constants to refer to specific objects of each type. For example, the subject 'Alice' is defined as *DefinitionAlice : subject* := 101. and the act 'Play' as *DefinitionPlay : act* := 301.. For each "nat" type in ODRL0 we have also used constants that play the role of "Null" objects (see "Null Object Pattern" [7]), for example *NullSubject*. This is needed partly because of the way ODRL0 language elements are defined which corresponds to the need to use *nonemptylist* exclusively even though at intermediate stages during the various algorithms Coq's *list* is a better fit because it allows empty lists.

Next we define the *policySet* datatype. Note the close/one-to-one mapping to its counterpart in listing 2.8. There are three ways a *policySet* can be constructed (see listing 2.8) corresponding to three constructors: *PrimitivePolicySet*, *PrimitiveExclusivePolicySet* and *AndPolicySet*. Both *PrimitivePolicySet* and *PrimitiveExclusivePolicySet* take a *preRequisite* and a *policy* as parameters. Finally *AndPolicySet* takes a non empty list of *policySet*s.

A *policy* is defined as a datatype with constructors *PrimitivePolicy* and *AndPolicy* (see listing 2.9). *PrimitivePolicy* takes a *preRequisite*, a *policyId* and a action *act*. Ignoring the *policyId* for a moment (it is only added to help the translation otherwise *policyId*s don't exist in ODRL proper), a primitive policy consists of a prerequisite and an action. If the prerequisite holds the action is allowed to be performed on the asset. The *AndPolicy* constructor is simply a *nonemptylist* (see listing 3.2) of *policy*s.

The data type *nonemptylist* reflects the definition of "policy conjunction" in listing listing 2.9 in chapter 2 (see definition of *nonemptylist* in listing 3.2). Essentially *nonemptylist* represents a list data structure that has at least one element and it is defined as a new *polymorphic* inductive type in its own section.

Listing 3.1: Coq version of agreement

```coq
Inductive agreement : Set :=
  | Agreement : prin → asset → policySet → agreement.

Definition prin := nonemptylist subject.

Definition asset := nat.

Definition subject := nat.

Definition act := nat.

Definition policyId := nat.

Inductive policySet : Set :=
  | PrimitivePolicySet : preRequisite → policy → policySet
  | PrimitiveExclusivePolicySet : preRequisite → policy → policySet
  | AndPolicySet : nonemptylist policySet → policySet.

Inductive policy : Set :=
  | PrimitivePolicy : preRequisite → policyId → act → policy
  | AndPolicy : nonemptylist policy → policy.
```

Listing 3.2: nonemptylist type

```coq
Section nonemptylist.

Variable X : Set.

Inductive nonemptylist : Set :=
  | Single : X → nonemptylist
  | NewList : X → nonemptylist → nonemptylist.

End nonemptylist.
```

In listing 3.3 *preRequisite* is defined as a new datatype with constructors *TruePrq*, *Constraint*, *ForEachMember*, *NotCons*, *AndPrqs*, *OrPrqs* and *XorPrqs* (see listing 2.12 for the abstract syntax equivalent).

TruePrq represents the always true prerequisite. The *Constraint* prerequisite is defined as the type *constraint* so its description is deferred here. Intuitively a constraint is a prerequisite to be satisfied that is outside the control of the user(s). For example, the constraint of being 'Alice' if you are 'Bob' (or 'Alice' for that matter). The constructor *ForEachMember* is defined to be a *prin* and a non empty list of *constraints*. Intuitively a *ForEachMember* prerequisite holds if each subject in *prin* satisfies each constraint in the list of *constraints*. The constructor *NotCons* is defined the same way the

12

Constraint"constructor is. This constructor is defined as the type *constraint* and it is meant to represent the negation of a *constraint* as we shall see in the translation (see listing 5.11). The remaining constructors *AndPrqs*, *OrPrqs* and *XorPrqs* take as parameters non empty lists of prerequisites. They represent conjunction, inclusive disjunction and exclusive disjunction of prerequisites respectively.

Finally a *constraint* (see listing 2.13 for the abstract syntax equivalent) is defined as a new datatype with constructors *Principal*, *Count* and *CountByPrin*.

Principal constraint takes a *prin* to match. For example, the constraint of the user being Bob would be represented as "Principal constraint". The *Count* constructor takes a *nat* which represents the number of times the user of an agreement has invoked the corresponding policies to justify her actions. If the count constraint is part of a policy then the corresponding policies is basically the single policy, whereas in the case that the count constraint is part of a policy set, the corresponding policies would be the set of those policies specified in the policy set. The *CountByPrin* is similar to *Count* but it takes an additional *prin* parameter. In this case the subjects specified in the *prin* parameter override the agreements' user(s).

Listing 3.3: preRequisite

```
Inductive preRequisite : Set :=
  | TruePrq : preRequisite
  | Constraint : constraint → preRequisite
  | ForEachMember : prin → nonemptylist constraint → preRequisite
  | NotCons : constraint → preRequisite
  | AndPrqs : nonemptylist preRequisite → preRequisite
  | OrPrqs : nonemptylist preRequisite → preRequisite
  | XorPrqs : nonemptylist preRequisite → preRequisite.


Inductive constraint : Set :=
  | Principal : prin → constraint
  | Count : nat → constraint
  | CountByPrin : prin → nat → constraint.
```

Chapter 4

ODRL0 Semantics

4.1 Introduction

In this section, we describe the semantics of ODRL0 language by a translation from agreements to a subset of many-sorted first-order logic formulas with equality. Note that in the listings in this chapter we use $[\![\,]\!]$ (double square brackets) notation as a mapping of ODRL0 syntactic elements to their translations as many-sorted first-order logic formulas. We also use $\triangleq$ between a translation and its corresponding formula to mean the translation is "mapped to" the formula.

The semantics will help answer queries of the form "may subject s perform action act to asset a?". If the answer is yes, we say permission is granted. Otherwise permission is denied.

At a high-level, an agreement is translated into a conjunction of formulas of the form $\forall x(prerequisites(x) \to P(x))$ where $P(x)$ itself is a conjunction of formulas of the form $prerequisites(x) \to (\neg)Permitted(x, act, a)$, where $Permitted(x, act, a)$ means the subject x is permitted to perform action act on asset a.

4.2 Agreement Translation

The translation of an *agreement* returns the translation for a *policySet* with arguments $prin_u$, the agreement's user and a, the asset.

Listing 4.1: Agreement Translation

$[\![agreement\ for\ prin_u\ about\ a\ with\ policySet]\!] \triangleq [\![policySet]\!]^{prin_u,a}$

14

4.3   Policy Set Translation

The translation for a *policySet* ($[\![policySet]\!]^{prin_u,a}$) is defined by cases, one for each clause of the grammar in listing 2.8. Recall that a *policySet* is either a *PrimitivePolicySet*, a *PrimitiveExclusivePolicySet* or a *AndPolicySet*. Each of these has its own translation function, which will be defined in the next 3 subsections.

4.3.1   PrimitivePolicySet Translation

Translation of a *PrimitivePolicySet* (*preRequisite → policy*) yields a formula that includes a test on whether the subject is in the set of agreements' users, the translation of the policy and the translation of the *prerequisite*. Basically if the subject in question is a user of the agreement and the policySet prerequisites hold, then the policy holds. Translation of the policy for a *PrimitivePolicySet* is called a *positive translation*. A positive translation is one where the actions described by the policies are permitted.

Listing 4.2: Policy Set Translation : PrimitivePolicySet

$$[\![preRequisite \to policy]\!]^{e,prin_u,a} \triangleq \forall x \ (([\![prin_u]\!]_x \ \wedge$$
$$[\![preRequisite]\!]_x^{e,getId(p),prin_u,a}) \to [\![policy]\!]_x^{positive,e,prin_u,a})$$

Listing 4.3: Positive Policy Translation : Single policy

$$[\![preRequisite \Rightarrow_{policyId} act]\!]_x^{positive,e,prin_u,a} \triangleq$$
$$([\![preRequisite]\!]_x^{e,policyId,prin_u}) \Rightarrow Permitted(x, [\![act]\!], a)$$

If the policy is a *AndPolicy*, the translation yields a conjunction of positive translations of each policy in turn.

Listing 4.4: Positive Policy Translation : List of policies

$$[\![and[policy_1, ..., policy_m]]\!]^{positive,e,prin_u,a} \triangleq [\![policy_1]\!]^{positive,e,prin_u,a} \ \wedge \ ... \ \wedge$$
$$[\![policy_m]\!]^{positive,e,prin_u,a}$$

4.3.2   PrimitiveExclusivePolicySet Translation

PrimitiveExclusivePolicySet (*preRequisite ↦ policy*) yields the conjunction of two implications. The first implication, is the same as one found in the translation of *PrimitivePolicySet*. The second implication however restricts access (to make the policy set exclusive) to only those subjects that are in the agreement's user. Translation of the policy in the second implication is called a *negative translation*. A negative translation is one where the actions described by the policies are not permitted.

Listing 4.5: Policy Set Translation : PrimitiveExclusivePolicySet

$$\llbracket preRequisite \mapsto policy \rrbracket^{e,prin_u,a} \triangleq \forall x \; ((\llbracket prin_u \rrbracket_x \; \wedge$$
$$\llbracket preRequisite \rrbracket_x^{e,getId(p),prin_u,a}) \to \llbracket policy \rrbracket_x^{positive,e,prin_u,a}) \; \wedge \; \forall x$$
$$(\neg \llbracket prin_u \rrbracket_x \to \llbracket policy \rrbracket_x^{negative,e,a})$$

Listing 4.6: Negative Policy Translation : Single policy

$$\llbracket preRequisite \Rightarrow_{policyId} act \rrbracket_x^{negative,e,prin_u,a} \triangleq$$
$$(\llbracket preRequisite \rrbracket_x^{e,policyId,prin_u}) \Rightarrow \neg(Permitted(x, \llbracket act \rrbracket, a))$$

If the policy is a *AndPolicy*, the translation yields a conjunction of negative translations of each policy in turn.

Listing 4.7: Negative Policy Translation : List of policies

$$\llbracket and[policy_1,...,policy_m] \rrbracket^{negative,e,a} \triangleq \llbracket policy_1 \rrbracket^{negative,e,a} \; \wedge \; ... \; \wedge \; \llbracket policy_m \rrbracket^{negative,e,a}$$

4.3.3 AndPolicySet Translation

AndPolicySet translates to conjunctions of the corresponding policy set translations.

Listing 4.8: Policy Set Translation : AndPolicySet

$$\llbracket and[policySet_1,...,policySet_m] \rrbracket^{e,prin_u,a} \triangleq \llbracket policySet_1 \rrbracket^{e,prin_u,a} \; \wedge \; ... \; \wedge$$
$$\llbracket policySet_m \rrbracket^{e,prin_u,a}$$

4.4 Principal Translation

Translation for a *prin* ($\llbracket prin \rrbracket_x$) is a formula that is true if and only if the subject x is in the prin set. A *prin* is either a single subject or a list of subjects ($\{subject_1,...,subject_m\}$ so the translation covers both cases. Each of these has its own translation function, which will be defined in the following 2 subsections.

If the *prin* is a single subject, the translation is a formula that is true if and only if the subject x is the same as the single subject *subject*.

4.4.1 Single Subject Translation

Listing 4.9: Prin Translation : Single subject

$$\llbracket subject \rrbracket_x \triangleq x = subject$$

16

4.4.2 List of Subjects Translation

Translation of a list of subjects is the disjunction of the translations for each subject.

Listing 4.10: Prin Translation : List of subjects

$$[\![\{subject_1, ..., subject_m\}]\!]_x \triangleq [\![subject_1]\!]_x \lor ... \lor [\![subject_m]\!]_x$$

4.5 Prerequisite Translation

Translation for a prerequisite is a formula $[\![prerequisite]\!]_x^{[id_1,...,id_m],prin,a}$, where the set of ids refer to identifiers for policies that are implied by the prerequisites, $prin$ is the agreement's user(s) (and to which the prerequisites apply), a is the asset and x is a variable of type $subject$. The translation for a $prerequisite$ is described by translation formulas for each type of $prerequisite$. A $prerequisite$ is either always $true$, a $Constraint$, a $ForEachMember$, a $NotCons$, a $AndPrqs$, a $OrPrqs$ or a $XorPrqs$. Each of these has its own translation function, which will be defined in the following subsections.

4.5.1 True Prerequisite Translation

The translation for a $TruePrq$ yields a formula that is always $true$.

Listing 4.11: Prerequisite Translation : Always True Prerequisite

$$[\![prerequisite :: true]\!] \triangleq \texttt{True}$$

4.5.2 Constraint Prerequisite Translation

The translation for a $Constraint$ is handled by a specialized constraint translation function (coverage of which starts at listing 4.18.

Listing 4.12: Prerequisite Translation : Constraint

$$[\![prerequisite :: constraint]\!]_x^{[id_1,...,id_m],prin_u} \triangleq [\![constraint]\!]_x^{[id_1,...,id_m],prin_u}$$

4.5.3 ForEachMember Prerequisite Translation

The translation for a $ForEachMember$ is also is handled by a specialized translation function (listing 4.21.

$$[\![prerequisite :: forEachMember]\!]_x^{[subject_1,...,subject_k],[constraint_1,...,constraint_m],[id_1,...,id_n]} \triangleq$$
$$[\![forEachMember]\!]_x^{[subject_1,...,subject_k],[constraint_1,...,constraint_m],[id_1,...,id_n]}$$

4.5.4 NotCons Prerequisite Translation

The translation for a *NotCons* yields a formula that is simply the negation of the translation for a constraint.

Listing 4.14: Prerequisite Translation : Not Constraint

$$[\![not\ prerequisite :: constraint]\!]_x^{[id_1,...,id_m],prin_u} \triangleq \neg [\![constraint]\!]_x^{[id_1,...,id_m],prin_u}$$

4.5.5 AndPrqs Prerequisite Translation

The translation for a *AndPrqs* yields a formula that is the conjunction of the translation for each *preRequisite*.

Listing 4.15: Prerequisite Translation : Conjunction

$$[\![and\ [preRequisite_1,...,preRequisite_k]]\!]^{[id_1,...,id_m],prin_u} \triangleq [\![preRequisite_1]\!]^{[id_1,...,id_m],prin_u} \wedge$$
$$... \wedge [\![preRequisite_k]\!]^{[id_1,...,id_m],prin_u}$$

4.5.6 OrPrqs Prerequisite Translation

The translation for a *OrPrqs* yields a formula that is the inclusive disjunction of the translation for each *preRequisite*.

Listing 4.16: Prerequisite Translation : Inclusive Disjunction

$$[\![or\ [preRequisite_1,...,preRequisite_k]]\!]^{[id_1,...,id_m],prin_u} \triangleq [\![preRequisite_1]\!]^{[id_1,...,id_m],prin_u} \vee$$
$$... \vee [\![preRequisite_k]\!]^{[id_1,...,id_m],prin_u}$$

4.5.7 XorPrqs Prerequisite Translation

The translation for a *XorPrqs* yields a formula that is the exclusive disjunction of the translation for each *preRequisite*.

Listing 4.17: Prerequisite Translation : Exclusive Disjunction

$$[\![Xor~[preRequisite_1,...,preRequisite_k]]\!]^{[id_1,...,id_m],prin_u} \triangleq [\![preRequisite_1]\!]^{[id_1,...,id_m],prin_u} \oplus$$
$$... \oplus [\![preRequisite_k]\!]^{[id_1,...,id_m],prin_u}$$

4.6   Constraint Translation

Translation for a constraint is a formula $[\![constraint]\!]_x^{[id_1,...,id_m],prin_u,a}$, where the set of ids refer to identifiers for policies that are implied by the constraint, $prin_u$ is the agreement's user(s) (and to which the constraint applies), a is the asset and x is a variable of type $subject$. The translation for a $constraint$ is described by translation formulas for each type of $constraint$. A $constraint$ is either a $Principal$, a $Count$, or a $CountByPrin$. Each of these has its own translation function, which will be defined in the following subsections.

4.6.1   Principal Constraint Translation

The translation for a $Principal$ is handled by a specialized translation function (listing 5.5.

Listing 4.18: Constraint Translation : Principal

$$[\![constraint :: prin]\!]_x^{[subject_1,...,subject_m]} \triangleq [\![prin]\!]_x^{[subject_1,...,subject_m]}$$

4.6.2   Count Constraint Translation

The translation for a $Count$ is handled by a specialized translation function (listing **??**.

Listing 4.19: Constraint Translation : Count

$$[\![constraint :: count[N]]\!]_x^{[id_1,...,id_m],prin_u} \triangleq [\![count[N]]\!]_x^{[id_1,...,id_m],prin_u}$$

4.6.3   CountByPrin Constraint Translation

The translation for a $CountByPrin$ is handled by the same specialized translation function as that for $Count$. The difference is that $CountByPrin$ overrides the subjects in $prin_u$ by a different set of subjects (listing **??**.

Listing 4.20: Constraint Translation : Count by Principal

$$[\![constraint :: prin(count[N])]\!]_x^{[subject_1,...,subject_m],[id_1,...,id_n]} \triangleq$$
$$[\![prin(count[N])]\!]_x^{[subject_1,...,subject_m],[id_1,...,id_n]}$$

4.6.4 forEachMember Translation

$$[\![forEachMember]\!]_x^{[subject_1,...,subject_k],[constraint_1,...,constraint_m],[id_1,...,id_n]} \triangleq$$
$$[\![constraint]\!]_x^{(subject_1,constraint_1),[id_1,...,id_n]} \wedge \ ... \ \wedge \ [\![constraint]\!]_x^{(subject_1,constraint_m),[id_1,...,id_n]}$$
$$\wedge \ ... \ \wedge \ [\![constraint]\!]_x^{(subject_2,constraint_1),[id_1,...,id_n]} \wedge \ ... \ \wedge$$
$$[\![constraint]\!]_x^{(subject_2,constraint_m),[id_1,...,id_n]} \wedge \ ... \ \wedge \ [\![constraint]\!]_x^{(subject_k,constraint_1),[id_1,...,id_n]}$$
$$\wedge \ ... \ \wedge \ [\![constraint]\!]_x^{(subject_k,constraint_m),[id_1,...,id_n]}$$

4.6.5 "Not Constraint" Translation

The translation for "Not Constraint" was listed in 4.14 earlier but we repeat it here to go along the Coq version.

$$[\![not \ \ constraint]\!]_x^{[id_1,...,id_m],prin_u} \triangleq \neg[\![constraint]\!]_x^{[id_1,...,id_m],prin_u}$$

4.7 Count Translation

Translation for *Count* or *CountByPrin* is based on whether the translation is on a single pair and multiple pairs of subject/policyIds. Each of these two cases will be described in the following subsections.

4.7.1 Count Translation For Subject/ID Pair

The translation for *Count* or *CountByPrin* for a pair of subject and policy identifier is a formula that is true if the number of times the $subject_1$ has invoked a policy with policy identifier id_1 is smaller than N.

$$[\![count[N]]\!]_x^{subject_1,id_1} \triangleq getCount(subject_1,id_1) < N$$

4.7.2 Count Translation For Subject/ID Pairs

The translation for *Count* or *CountByPrin* for subject and policy identifier pairs is a formula that is true if the total number of times that a subject has invoked a policy with policy identifier id_i is smaller than N.

Chapter 5

ODRL0 Semantics In Coq

5.1 Introduction

The translation functions plus the auxiliary types and infrastructure, implementing the semantics have been encoded in Coq. Translation functions build Coq terms of type *Prop*. Well-formed propositions (or *Prop*s) are assertions one can express about values such as mathematical objects or even programs (e.g. $3 < 8$) in Coq. Assertions may be true, false or simply conjectures, however an assertion is only true in general if a proof is provided ([2]).

Whether a permission is granted or denied depends on the agreements in question but also on the facts recorded in the environment. For ODRL0 those facts revolve around the number of times a policy has been used to justify an action (see section 2.2 for more details on odrl0). We encode this information in an *environment* which is a conjunction of equalities of the form $count(s, policyId) = n$.

The Coq version of the count equality is a new inductive type called *count_ equality*. An environment is defined to be a non-empty list of *count_equality* objects (see listing 5.1).

Listing 5.1: Environments and Counts

```
Inductive count_equality : Set :=
  | CountEquality : subject → policyId → nat → count_equality.


Inductive environment : Set :=
  | SingleEnv : count_equality → environment
  | ConsEnv : count_equality →  environment → environment.
```

We also define a *getCount* function (see listing 5.2) that given a pair consisting of a subject and policy id, looks for a corresponding count in the environment. *getCount* assumes the given environment is consistent, so it returns the first matched *count* it sees for a (*subject, id*) pair. If a *count* for a (*subject, id*) pair is not found it returns 0.

Listing 5.2: getCount Function

```
Fixpoint getCount
 (e:environment)(s:subject)(id: policyId): nat :=
  match e with
  | SingleEnv f ⇒
      match f with
      | CountEquality s1 id1 n1 ⇒
          if (beq_nat s s1)
          then if (beq_nat id id1) then n1 else 0
          else 0
      end
  | ConsEnv f rest ⇒
      match f with
      | CountEquality s1 id1 n1 ⇒
          if (beq_nat s s1)
          then if (beq_nat id id1) then n1 else (getCount rest s id)
          else (getCount rest s id)
      end
  end.
```

5.2 Translations

Translation of the top level *agreement* element proceeds by case analysis on the structure
of the agreement. However an agreement can only be built one way; by calling the con-
structor *Agreement*. The translation proceeds by calling the translation function for the
corresponding *policySet* namely the parameter to *Agreement* called *ps*.

Listing 5.3: Translation of agreement

```
Definition trans_agreement (e:environment)(ag:agreement) : Prop :=
 match ag with
   | Agreement prin_u a ps ⇒ trans_ps e ps prin_u a
 end.
```

Translation of a *policySet* (called $trans\_ps$ in listing 5.4), takes as input e, the environ-
ment, ps, the policy set, $prin_u$, the agreement's user, and a, the asset, and proceeds by case
analysis of different policySet constructors and recursing into translation functions for the
composing elements. A policySet is either a $PrimitivePolicySet$, $PrimitiveExclusivePolicySet$
or a $AndPolicySet$.

Note that to implement the translation for an $AndPolicySet$ a local function $trans_ps_list$
has been defined where for a single *policySet*, $trans_ps$ is called, and for a list of *policySet*s,
the conjunction of $trans_ps$s are returned.

Listing 5.4: Translation of Policy Set

```
Fixpoint trans_ps
  (e:environment)(ps:policySet)(prin_u:prin)(a:asset){struct ps} : Prop :=

let trans_ps_list := (fix trans_ps_list (ps_list:nonemptylist policySet)(prin_u:prin)
      (a:asset){struct ps_list}:=
  match ps_list with
    | Single ps1 ⇒ trans_ps e ps1 prin_u a
    | NewList ps ps_list' ⇒ ((trans_ps e ps prin_u a) /\ (trans_ps_list ps_list' prin_u a
      ))
  end) in
    match ps with
    | PrimitivePolicySet prq p ⇒ ∀ x, (((trans_prin x prin_u) /\
                          (trans_preRequisite e x prq (getId p) prin_u)) →
                          (trans_policy_positive e x p prin_u a))

    | PrimitiveExclusivePolicySet prq p ⇒ ∀ x, ((((trans_prin x prin_u) /\
                                (trans_preRequisite e x prq (getId p) prin_u)) →
                                (trans_policy_positive e x p prin_u a)) /\
                                ((not (trans_prin x prin_u)) → (trans_policy_negative e
      x p a)))

    | AndPolicySet ps_list ⇒ trans_ps_list ps_list prin_u a
    end.
```

Translation of a *prin* (called *trans_prin* in listing 5.5) takes as input x, the *subject* in question, p, the principal or the *prin*, and proceeds based on whether p is a single subject or a list of subjects. If p is a single subject, s, the *Prop* $x = s$ is returned. Otherwise the disjunction of the translation of the first subject in p (s) and the *rest* of the subjects is returned.

Listing 5.5: Translation of a Prin

```
Fixpoint trans_prin
  (x:subject)(p: prin): Prop :=

  match p with
    | Single s ⇒ (x=s)
    | NewList s rest ⇒ ((x=s) \/ trans_prin x rest)
  end.
```

A positive translation for a policy (called *trans_policy_positive* in listing 5.6) takes as input e, the *environment*, x, the *subject*, p, the *policy* to translate, $prin_u$, the agreement's user, and a, the asset and proceeds based on whether we have a *PrimitivePolicy* or a *AndPolicy*. If the policy is a *PrimitivePolicy* an implication is returned which indicates x is *permitted* to do *action* to a, if the *preRequisite* holds.

Permitted is a predicate specified as *ParameterPermitted* : *subject*− > *act*− > *asset*− > *Prop*. So *Permitted* predicate takes a *subject*, an *act* (an action) and an *asset* and builds a term of type *Prop*.

Note that to implement the translation for an *AndPolicy* a local function *trans_p_list* has been defined where for a single *policy*, *trans_policy_positive* is returned, and for a list of *policy*s, the conjunction of *trans_policy_positive*s are returned.

Listing 5.6: Translation of a positive policy

```
Fixpoint trans_policy_positive
  (e:environment)(x:subject)(p:policy)(prin_u:prin)(a:asset){struct p} : Prop :=

let trans_p_list := (fix trans_p_list (p_list:nonemptylist policy)(prin_u:prin)(a:
    asset){struct p_list}:=
              match p_list with
                | Single p1 ⇒ trans_policy_positive e x p1 prin_u a
                | NewList p p_list' ⇒
                    ((trans_policy_positive e x p prin_u a) /\
                     (trans_p_list p_list' prin_u a))
              end) in


  match p with
    | PrimitivePolicy prq policyId action ⇒ ((trans_preRequisite e x prq (Single
    policyId) prin_u) →
                                      (Permitted x action a))
    | AndPolicy p_list ⇒ trans_p_list p_list prin_u a
  end.
```

A negative translation for a policy (called *trans_policy_negative* in listing 5.7) takes as input *e*, the *environment*, *x*, the *subject*, *p*, the *policy* to translate, and *a* the asset and proceeds based on whether we have a *PrimitivePolicy* or a *AndPolicy*. If the policy is a *PrimitivePolicy* an implication is returned which indicates *x* is forbidden to do *action* to *a* regardless of whether *preRequisite* holds. Note that the notation (¬) indicates that *Permitted* may be negated. As the case for the positive translation, to implement the translation for an *AndPolicy* a local function *trans_p_list* has been defined where for a single *policy*, *trans_policy_negative* is returned, and for a list of *policy*s, the conjunction of *trans_policy_negative*s are returned.

```
Fixpoint trans_policy_negative
  (e:environment)(x:subject)(p:policy)(a:asset){struct p} : Prop :=
let trans_p_list := (fix trans_p_list (p_list:nonemptylist policy)(a:asset){
    struct p_list}:=
            match p_list with
              | Single p1 ⇒ trans_policy_negative e x p1 a
              | NewList p p_list' ⇒ ((trans_policy_negative e x p a) /\
                              (trans_p_list p_list' a))
            end) in


  match p with
    | PrimitivePolicy prq policyId action ⇒ not (Permitted x action a)
    | AndPolicy p_list ⇒ trans_p_list p_list a
  end.
```

The translation of a *prerequisite* (called *trans_preRequisite* in listing 5.8) takes as input *e*, the *environment*, *x*, the *subject*, *prq*, the *preRequisite* to translate, *IDs*, the set of identifiers (of policies implied by the *prq*), *prin_u*, the agreement's user, and proceeds by case analysis on the structure of the *prerequisite*. A *prerequisite* is either a *TruePrq*, a *Constraint*, a *ForEachMember*, a *NotCons*, a *AndPrqs*, a *OrPrqs* or a *XorPrqs*.

In listing 5.8 the translation for *TruePrq* is the Prop *True*, the translations for *Constraint*, *ForEachMember* and *NotCons* simply call respective translation functions for corresponding types *constraint* and *forEachMember* (namely *trans_constraint*, *trans_forEachMember* and *trans_notCons*). Note that the translation for *AndPrqs*, *OrPrqs* and *XorPrqs* have not yet been implemented but based on the their many-sorted-logic formulas' specifications (listing 2.12) they will be conjunctions, disjunctions and exclusive disjunctions of translations for each *prerequisite*.

Listing 5.8: Translation of a PreRequisite

```
Definition trans_preRequisite
  (e:environment)(x:subject)(prq:preRequisite)(IDs:nonemptylist policyId)(prin_u:prin) :
    Prop :=

  match prq with
    | TruePrq ⇒ True
    | Constraint const ⇒ trans_constraint e x const IDs prin_u
    | ForEachMember prn const_list ⇒ trans_forEachMember e x prn const_list IDs
    | NotCons const ⇒ trans_notCons e x const IDs prin_u
    | AndPrqs prqs ⇒ True
    | OrPrqs prqs ⇒ True
    | XorPrqs prqs ⇒ True
  end.
```

The translation of a *constraint* (called *trans_constraint* in listing 5.9) takes as input *e* the *environment*, *x* the *subject*, *const*, the *constraint* to translate, *IDs*, the set of identifiers (of policies implied by the parent *preRequisite*) and $prin_u$, the agreement's user and proceeds by case analysis on the structure of the *constraint*. A *constraint* is either a *Principal*, a *Count* or a *CountByPrin*.

In listing 5.9 the translation for *Principal* returns the translation function (namely *trans_prin*) for the *prn* (the *prin* that accompanies the *const* constraint). The translation for *Count* and *CountByPrin* return the translation function *trans_count*. For *Count* the *prin* used is the agreement's user, whereas the *prin* used is the one passed to *CountByPrin* namely *prn*.

Listing 5.9: Translation of a Constraint

```
Fixpoint trans_constraint
  (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)
  (prin_u:prin){struct const} : Prop :=
  match const with
    | Principal prn ⇒ trans_prin x prn

    | Count n ⇒ trans_count e n IDs prin_u

    | CountByPrin prn n ⇒ trans_count e n IDs prn

  end.
```

The translation of a *forEachMember* (called *trans_forEachMember* in listing 5.10) takes as input *e* the *environment*, *x* the *subject*, *principals*, the set of subjects that override the agreement's user(s), *const_list* the set of constraints and *IDs*, the set of identifiers (of policies implied by the parent *preRequisite*).

To implement the translation for a *forEachMember* we start by calling an auxiliary function *process_two_lists* that effectively returns a new list composed of pairs of members of the first list and the second list (the cross-product of the two input lists). In the case of a *forEachMember* translation, the call is "*process_two_lists principals const_list*" which returns a list of pairs of subject and constraint namely *prins_and_constraints*. *prins_and_constraints* is then passed to a locally defined function *trans_forEachMember_-Aux* where for a single pair of subject and constraint *trans_constraint* is called and for a list of pairs of subject and constraints, the conjunction of *trans_constraint*s (for the first pair) and *trans_forEachMember_Aux*s (for the rest of the pairs) are returned.

Listing 5.10: Translation of forEachMember

```
Fixpoint trans_forEachMember
        (e:environment)(x:subject)(principals: nonemptylist subject)(const_list:
    nonemptylist constraint)
        (IDs:nonemptylist policyId){struct const_list} : Prop :=


let trans_forEachMember_Aux
  := (fix trans_forEachMember_Aux
        (prins_and_constraints : nonemptylist (Twos subject constraint))
        (IDs:nonemptylist policyId){struct prins_and_constraints} : Prop :=

    match prins_and_constraints with
      | Single pair1 ⇒ trans_constraint e x (right pair1) IDs (Single (left pair1))
      | NewList pair1 rest_pairs ⇒
          (trans_constraint e x (right pair1) IDs (Single (left pair1))) /\
          (trans_forEachMember_Aux rest_pairs IDs)
    end) in

    let prins_and_constraints := process_two_lists principals const_list in
    trans_forEachMember_Aux prins_and_constraints IDs.
```

The translation of a *NotCons* (called *trans_notCons* in listing 5.11) takes as input *e* the *environment*, *x* the *subject*, *const*, the *constraint* to translate, *IDs*, the set of identifiers (of policies implied by the parent *preRequisite*) and $prin_u$, the agreement's user and proceeds to return the negation of *trans_constraint* (see listing 5.9).

Listing 5.11: Translation of not cons

```
Definition trans_notCons
  (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:prin) :
    Prop :=
  ~ (trans_constraint e x const IDs prin_u).
```

The translation of a *Count* or a *CountByPrin* (called *trans_count* in listing 5.12) takes as input *e* the *environment*, *n* the total number of times the subjects mentioned in $prin_u$ (last parameter) may invoke the policies identified by *IDs* (third parameter).

To implement the translation for a *Count* or a *CountByPrin* we start by calling an auxiliary function *process_two_lists* that effectively returns a new list composed of pairs of members of the first list and the second list (the cross-product of the two input lists). In the case of *trans_count*, the call is "*process_two_lists IDs prin_u*" which returns a list of pairs of *policyId* and *subject* namely *ids_and_subjects*. *ids_and_subjects* is then passed to a locally defined function *trans_ count_ aux*.

trans_count_aux returns the current count for a single pair of *policyId* and *subject* (the call to *getCount* which looks up the environment *e* and returns the current count per

each *subject* and *policyId*) and for a list of pairs of *policyId* and *subject*s, the addition of *get_count* (for the first pair) and *trans_count_aux*s (for the rest of the pairs) is returned.

A local variable $running_total$ has the value returned by *trans_count_aux*. Finally the proposition $running_total < n$ is returned as the translation for a *Count* or a *CountByPrin*.

Note that the only difference between translations for a *Count* and a *CountByPrin* is the additional *prn* parameter for *CountByPrin* which allows for getting counts for subjects not necessarily the same as $prin_u$, the agreement's user(s).

Listing 5.12: Translation of count

```
Fixpoint trans_count
  (e:environment)(n:nat)(IDs:nonemptylist policyId)
  (prin_u:prin) : Prop :=

  let trans_count_aux
    := (fix trans_count_aux
         (ids_and_subjects : nonemptylist (Twos policyId subject)) : nat :=
      match ids_and_subjects with
        | Single pair1 ⇒ getCount e (right pair1) (left pair1)
        | NewList pair1 rest_pairs ⇒
            (getCount e (right pair1)(left pair1)) +
            (trans_count_aux rest_pairs)
      end) in

  let ids_and_subjects := process_two_lists IDs prin_u in
  let running_total := trans_count_aux ids_and_subjects in
  running_total < n.
```

Chapter 6

Queries

6.1 Introduction

We first mentioned queries in chapter 4 on page 14. Ultimately policy statements describing an agreement will be used to enforce those agreements. To enforce policy agreements, access queries are asked from the policy engine and access is granted or denied based on the answer.

By defining formal semantics for ODRL the authors of [8] were able to prove that answering a query on whether access should be granted or not, is decidable and NP-hard for the full ODRL.

In this chapter we will review our encoding of queries in Coq and Coq representations of other definitions used in [8] which we will use to prove decidability results of our own.

6.2 Queries

Queries are tuples of the form $(A, s, action, a, e)$ in [8]. The tuple corresponds to the question of determining whether a set A of agreements imply that a subject s may perform action $action$ on an asset a given the environment e. The Coq representation is listed in listing 6.1. We distinguish single agreement queries from multiple agreement queries by defining two separate types: $single\_query$ and $general\_query$.

Listing 6.1: Queries

```
Inductive single_query : Set :=
  | SingletonQuery : agreement → subject → act → asset → environment →
    single_query.


Inductive general_query : Set :=
  | GeneralQuery : nonemptylist agreement → subject → act → asset →
    environment → general_query.
```

6.3   Answering Queries

Answering a query as defined earlier can lead to one of four outcomes: error(listing 6.2), permitted(listing 6.3), denied(listing 6.4) and "not applicable" (listing 6.5) defined in [9].

Listing 6.2: Answerable Queries: Error

$$(\bigwedge[\![agreement]\!]) \wedge E \implies Permitted(s, act, a) \text{ and}$$
$$(\bigwedge[\![agreement]\!]) \wedge E \implies \neg Permitted(s, act, a)$$

Listing 6.3: Answerable Queries: Permit

$$(\bigwedge[\![agreement]\!]) \wedge E \implies Permitted(s, act, a) \text{ and}$$
$$(\bigwedge[\![agreement]\!]) \wedge E \implies\!\!\!\!\!/ \ \neg Permitted(s, act, a)$$

Listing 6.4: Answerable Queries: Deny

$$(\bigwedge[\![agreement]\!]) \wedge E \implies\!\!\!\!\!/ \ Permitted(s, act, a) \text{ and}$$
$$(\bigwedge[\![agreement]\!]) \wedge E \implies \neg Permitted(s, act, a)$$

Listing 6.5: Answerable Queries: Not Applicable

$$(\bigwedge[\![agreement]\!]) \wedge E \implies\!\!\!\!\!/ \ Permitted(s, act, a) \text{ and}$$
$$(\bigwedge[\![agreement]\!]) \wedge E \implies\!\!\!\!\!/ \ \neg Permitted(s, act, a)$$

In [8] a slightly different formulation is used to denote the same four decision types. "Query Inconsistent","Permission Granted", "Permission Denied" and "Permission Unregulated". Define the formulas f_q^+ and f_q^- (see listings 6.6 and 6.7).

Listing 6.6: f_q^+

$$f_q^+ \triangleq (\bigwedge[\![agreement]\!]) \implies Permitted(s, act, a)$$

30

Listing 6.7: f_q^-

$$f_q^- \triangleq \left(\bigwedge \llbracket agreement \rrbracket \right) \implies \neg Permitted(s, act, a)$$

Now answering the queries will depend on the *E-validity* of f_q^+ and f_q^-. E-validity or the consistency of the environment is not captured explicitly in [8] (see listings 6.8, 6.9, 6.10 and 6.11) however the decision of which answer a query results in takes the consistency of the environment into account. In this thesis we will encode and use the decision algorithms in [8].

Listing 6.8: Answerable Queries: Query Inconsistent

```
f_q^+ and f_q^- both hold
```

Listing 6.9: Answerable Queries: Permission Granted

```
f_q^+ holds and f_q^- does not hold
```

Listing 6.10: Answerable Queries: Permission Denied

```
f_q^+ does not hold and f_q^- holds
```

Listing 6.11: Answerable Queries: Permission Unregulated

```
f_q^+ does not hold and f_q^- does not hold
```

Chapter 7

Examples

7.1 Introduction

In this chapter we will take a tour of the syntax and semantics we have so far developed by examining some example agreements. In the following we will start by reviewing some of the examples used in [8].

Ultimately the goal of specifying all the syntax and semantics is to declare some interesting theorems about policy expressions and proving them in Coq however we will start with some specific propositions/theorems about these examples to get a feel for how proofs are done in Coq.

7.2 Agreement 2.1

Consider example 2.1 (from [8]) where the *policySet* is a *AndPolicySet* with $p1$ and $p2$ as the individual *policySet*s. Let $p1$ be defined as $Count[5] \rightarrow print$ and $p2$ as $and[Alice, Count[2]] \rightarrow print$.

The agreement is that the asset *The Report* may be printed a total of five times by either *Alice* or *Bob*, and twice more by Alice. So if Alice and Bob have used policy $p1$ to justify their printing of the asset m_{p1} and n_{p1} times, respectively, then either may do so again if $m_{p1} + n_{p1} < 5$. If they have used $p2$ to justify their printing of the asset m_{p2} and n_{p2} times, respectively, then only Alice may do so again if $m_{p2} + n_{p2} < 2$. Note that since Bob doesn't meet the prerequisite of being Alice, n_{p2} is effectively 0, so we have $m_{p2} < 2$ as the condition for Alice being able to print again (Alice does meet the prerequisite of being Alice).

Listing 7.1: Agreement 2.1 (as used in [8])

```
agreement
 for {Alice, Bob}
 about TheReport
 with and [p1, p2].
```

The Coq version of the agreement 2.1 (listing 7.1) and its sub-parts is listed below. It is best to start with the agreement itself called *A2.1* in the listing and compare to the agreement 2.1 listed in 7.1.

Listing 7.2: Agreement 2.1 in Coq

```
Definition p1A1:policySet :=
  PrimitivePolicySet
    TruePrq
    (PrimitivePolicy (Constraint (Count  5)) id1 Print).

Definition p2A1prq1:preRequisite := (Constraint (Principal (Single Alice))).
Definition p2A1prq2:preRequisite := (Constraint (Count 2)).

Definition p2A1:policySet :=
  PrimitivePolicySet
    TruePrq
    (PrimitivePolicy (AndPrqs (NewList p2A1prq1 (Single p2A1prq2))) id2 Print).

Definition A2.1 := Agreement (NewList Alice (Single Bob)) TheReport
                (AndPolicySet (NewList p1A1 (Single p2A1))).
```

7.3   Agreement 2.5

Consider example 2.5 (from [8]) where the *policySet* is a *PrimitivePolicySet* with a *Count* constraint as prerequisite and a *AndPolicy* as the policy. The *AndPolicy* is the conjunction of two *PrimitivePolicy*s. Both policies have prerequisites of type *ForEachMember* with actions *display* and *print* respectively. The *prin* component for both *ForEachMember*s is *Alice, Bob*, whereas the constraint for the first *ForEachMember* is *Count*[5] and for the second is *Count*[2].

Listing 7.3: Agreement 2.5 (as used in [8])

```
agreement
 for {Alice, Bob}
 about ebook
 with Count [10] → and [forEachMember[{Alice, Bob}; Count[5]] ⇒_id1 display,
                        forEachMember[{Alice, Bob}; Count[1]] ⇒_id2 print].
```

The Coq version of the agreement 2.5 (listing 7.3) and its sub-parts is listed below. See agreement 2.5 listed in 7.3 for comparison.

The agreement is that the asset *ebook* may be displayed up to five times by Alice and Bob each, and printed once by each. However the total number of actions (either *display* or *print*) justified by the two policies by either Alice and Bob is at most 10.

Listing 7.4: Example 2.5

```
Definition tenCount:preRequisite := (Constraint (Count 10)).
Definition fiveCount:constraint := (Count 5).
Definition oneCount:constraint := (Count 1).

Definition prins2_5 := (NewList Alice (Single Bob)).
Definition forEach_display:preRequisite := ForEachMember prins2_5 (Single fiveCount)
    .
Definition forEach_print:preRequisite := ForEachMember prins2_5 (Single oneCount).

Definition primPolicy1:policy := PrimitivePolicy forEach_display id1 Display.
Definition primPolicy2:policy := PrimitivePolicy forEach_print id2 Print.

Definition policySet2_5:policySet :=
  PrimitivePolicySet tenCount (AndPolicy (NewList primPolicy1 (Single primPolicy2))).

Definition A2.5 := Agreement prins2_5 ebook policySet2_5.
```

7.4   Agreement 2.6

Consider example 2.6 (from [8]) where the *policySet* is a *PrimitiveExclusivePolicySet*
with a *InSequence requirement* as prerequisite. We will cover requirements type con-
straints in *ODRL1* since in ODRL0 we have elided their use. We will also describe this
example in detail when ODRL1 constructs are added.

Listing 7.5: Example 2.6

```
Definition prins2_6 := prins2_5.

Definition aliceCount10:preRequisite := Constraint (CountByPrin (Single Alice) 10).
Definition primPolicy2_6:policy := PrimitivePolicy aliceCount10 id3 Play.
Definition policySet2_6_modified:= PrimitiveExclusivePolicySet TruePrq
    primPolicy2_6.
```

Chapter 8

Some Simple Theorems

8.1 Introduction

In this chapter we will declare and prove some very simple theorems about the examples from chapter [7]. This simple introduction is only meant to give us a feel for how theorems are stated in Coq and how proofs are constructed using Coq *tactics*.

Propositions are types in Coq whose type is the sort *Prop*. Any term t whose type is a proposition is a proof term or, for short, a proof. A *Hypothesis* is a local declaration $h : P$ where h is an identifier and P is a proposition. An *Axiom* is similar to a hypothesis except it is declared at the global scope and so it is always available. A *Theorem* or *Lemma* is any identifier whose type is a proposition ([2]). Keywords "Hypothesis", "Axiom" and "Theorem" or "Lemma" are used in each case respectively.

To build a proof in Coq the user states the proposition to prove; this is called a goal to be proved or discharge, along with some hypothesis that makes up the local context. The user then uses commands called tactics to manipulate the local context and to decompose the goal into simpler goals. The goal simplification into sub-goals will continue until all the sub-goals are solved.

In listing 8.1 we have declared a theorem called $example1$ and the corresponding proposition $forall x : nat, x < x + 1$.

Note that the notation $P : T$ is also used to declare program P has type T. This duality of notation is due to Curry-Howard isomorphism which relates the two worlds of type theory and structural logic together. Once the Theorem has been declared Coq displays the proposition to be proved under a horizontal line written ———, and displays the context of local facts and hypothesis, if any, above the horizontal line. At this point one can enter proof mode by using *Proof.* upon which Coq is ready to accept tactics. Entering tactics that can break the stated goal (under the horizontal line) into one or more sub-goals is how one progresses until no goals left at which point Coq responds with "No more subgoals" ([1]).

Listing 8.1: Proof Example

```
Theorem example1: ∀ x:nat, x < x + 1.
```

8.2   Theorem One

In listing 8.3 we define a *policySet* with a *constraint* such that if *Alice* has used the policy with *id*1 to justify her printing a_1 times, she may do so again if $a_1 < 5$. The agreement *AgreeCan* simply links the asset *TheReport* with the subject *Alice* and the *policySet* previously defined.

We capture the fact that *Alice* has used the policy with *id*1 to justify her printing 2 times in an environment called *eA*1. Recall that environments are defined to be non-empty lists of $count_e quality$ objects (see listing 5.1).

We also declare a hypothesis *H* with the proposition that results from the translation of the agreement (see definition of $trans_a greement$ in listing 5.3) and the environment. The proposition can be shown in Coq after some clean-up (e.g. replaced 101 by Alice) and using the form *Eval compute* :

Listing 8.2: Hypothesis for Theorem One
```
∀ x : subject, x = Alice /\ True → 2 < 5 → Permitted x Print TheReport.
```

The theorem *One* that we are going to prove is trivial but nonetheless in English it states that Alice is Permitted to Print TheReport. The proof comes after the command 'Proof.' and ends with 'Qed'.

Listing 8.3: Theorem One
```
Definition psA1:policySet :=
 PrimitivePolicySet
   TruePrq
   (PrimitivePolicy (Constraint (Count  5)) id1 Print).

Definition AgreeCan := Agreement (Single Alice) TheReport psA1.

Definition eA1 : environment :=
  (SingleEnv (make_count_equality Alice id1 2)).

Hypothesis H: trans_agreement eA1 AgreeCan.

Theorem One: Permitted Alice Print TheReport.

Proof. simpl in H. apply H. split. reflexivity. auto.  omega. Qed.
```

8.3 Theorem Two

In listing 8.5 we define an exclusive policy set *policySet* containing a policy *pol* that allows printing. The agreement *AgreeA5* includes the exclusive policy set to express that Bob may print *LoveAndPeace*. However any subject that is not the agreement's user (e.g. Bob) is forbidden from printing *LoveAndPeace*.

Notice that due to the fact that environments are defined as non-empty lists, we have added a Null count to it (see *eA5*). We continue to capture the relevant facts from the environment and the agreement through defining a hypothesis (e.g. *H*). The hypothesis is shown below :

Listing 8.4: Hypothesis for Theorem Two

```
∀ x : subject,
     (x = Bob /\ True → True → Permitted x Print LoveAndPeace) /\
     ((x = Bob → False) → Permitted x Print LoveAndPeace → False).
```

Theorem $T1\_A5$ states the exclusivity of the policy set, namely that any subject that is not Bob is not permitted to print the asset LoveAndPeace. Theorem $T2\_A5$ uses $T1\_A5$ to prove Alice is not permitted to print the asset.

Listing 8.5: Theorem Two

```
Definition prin_bob := (Single Bob).
Definition pol:policy := PrimitivePolicy TruePrq id3 Print.
Definition pol_set:policySet := PrimitiveExclusivePolicySet TruePrq pol.
Definition AgreeA5 := Agreement prin_bob LoveAndPeace pol_set.
Definition eA5 : environment := (SingleEnv (make_count_equality NullSubject NullId 0))
   .


Hypothesis H: trans_agreement eA5 AgreeA5.


Theorem T1_A5: ∀ x, x<>Bob →  ~Permitted x Print LoveAndPeace.
Proof. simpl in H. apply H. Qed.

Theorem T2_A5: ~Permitted Alice Print LoveAndPeace.
Proof. simpl in H. apply T1_A5. apply not_eq_S. omega. Qed.


End A5.
```

8.4 Theorem Three

In listing 5.1 we defined environments as non-empty lists of $count_equality$ objects which are in turn defined as counts per each subject, policy-id pair. These count formulas represent how many times each policy has been used to justify an action by a subject wrt a policy (specified by the policy id) and semantically it makes sense that they are unique in time. When and if two $count_equality$ objects with the same subject and policy id refer to different counts, we say we have *inconsistent* count formulas. The listing 8.6 defines the binary predicate *inconsistent*.

Listing 8.6: Inconsistent Count Formulas

```
Definition inconsistent (f1 f2 : count_equality) : Prop :=
  match f1 with (CountEquality s1 id1 n1) ⇒
    match f2 with (CountEquality s2 id2 n2) ⇒
      s1 = s2 → id1 = id2 → n1 <> n2
    end
  end.
```

Next we would like to expand the notion of inconsistency to more than two count formulas. We first define a predicate over a count formula and an environment as in listing 8.7. If the environment is a singleton then we just compare the two count formulas for inconsistency, else we build the disjunction of the inconsistency between the count formula on one hand and the head of the environment and the rest of the environment, respectively.

Listing 8.7: Inconsistent Count Formula And Environment

```
Fixpoint formula_inconsistent_with_env (f: count_equality)
                    (e : environment) : Prop :=
 match e with
   | SingleEnv g ⇒ inconsistent f g
   | ConsEnv g rest ⇒ (inconsistent f g) \/ (formula_inconsistent_with_env f rest)
 end.
```

Finally we define a new inductive data type that represents *consistent* environments (see listing 8.8). An environment is consistent if it is a singleton count formula, if it consists of only two consistent count formulas and finally if the environment consists of a consistent environment and the consistent composition of a count formula and the consistent environment (see constructor *consis_ more*.

Listing 8.8: Inconsistent Environment

```
Inductive env_consistent : environment → Prop :=
| consis_1 : ∀ f, env_consistent (SingleEnv f)
| consis_2 : ∀ f g, ~(inconsistent f g) → env_consistent (ConsEnv f (SingleEnv g))
| consis_more : ∀ f e,
```

```
env_consistent e → ~(formula_inconsistent_with_env f e) → env_consistent (
  ConsEnv f e).
```

We will now pose several small theorems about consistency of count formulas and environments and provide proofs for them (see listing 8.9).

Listing 8.9: Inconsistent Environment

```
Theorem f1_and_f2_are_inconsistent: inconsistent f1 f2.
Proof.
unfold inconsistent. simpl. omega. Qed.


Theorem f1_and___env_of_f2_inconsistent: formula_inconsistent_with_env f1 (
    SingleEnv f2).
Proof.
unfold formula_inconsistent_with_env. apply f1_and_f2_are_inconsistent. Qed.


Theorem two_inconsistent_formulas_imply_env_inconsistent:
 ∀ f g, inconsistent f g → ~env_consistent (ConsEnv f (SingleEnv g)).
Proof.
intros. unfold not. intros H'.
inversion H'. intuition. intuition. Qed.


Theorem e2_is_inconsistent: ~env_consistent e2.
Proof.
apply two_inconsistent_formulas_imply_env_inconsistent.
apply f1_and_f2_are_inconsistent. Qed.


Theorem env_consistent_implies_two_consistent_formulas:
 ∀ (f g: count_equality),
    env_consistent (ConsEnv f (SingleEnv g))→ ~inconsistent f g.
Proof.
intros. inversion H. exact H1. intuition. Qed.


Theorem two_consistent_formulas_imply_env_consistent:
 ∀ (f g: count_equality),
    ~inconsistent f g → env_consistent (ConsEnv f (SingleEnv g)).
Proof.
intros. apply consis_2. exact H. Qed.

Theorem env_inconsistent_implies_two_inconsistent_formulas:
 ∀ (f g: count_equality),
    ~env_consistent (ConsEnv f (SingleEnv g))→ inconsistent f g.
```

```
Proof.
induction f.
induction g.
unfold inconsistent.
intros.
subst.
generalize (dec_eq_nat n n0).
intro h; elim h.
intro; subst.
elim H.
apply consis_2.
unfold inconsistent.
intro.
assert (s0=s0); auto.
assert (p0=p0); auto.
specialize H0 with (1:=H1) (2:=H2).
elim H0; auto.
auto.
Qed.


Theorem same_subjects_policyids_different_counts_means_inconsistent : ∀ (s1 s2:
    subject),
            ∀ (id1 id2: policyId),
            ∀ (n1 n2: nat),

  (s1 = s2 /\ id1 = id2 /\ n1 <> n2) →
  inconsistent (CountEquality s1 id1 n1) (CountEquality s2 id2 n2).

Proof.
intros. unfold inconsistent. intros. intuition. Qed.
```

Chapter 9

Proposed Future Work

9.1 Introduction

9.2 Machine-Checked Proof of Decidability of Queries

By defining formal semantics for ODRL authors of [8] were able to show some important results. First result is that answering the question of whether a set of ODRL statements imply a permission, denial or other possibilities is decidable and also that its complexity is NP-hard (see Theorem 4.1 from [8] re-printed below).

Theorem 4.1 *The problem of deciding, for a query $q = (A, s, act, a, E)$, whether f_q^+ is E-valid is decidable but NP-hard. Similarly for f_q^-.*

The authors of [8] then prove that by removing the construct *not*[*policySet*] from ODRL's syntax answering the same query remains decidable and efficient (polynomial time complexity).

We will prove equivalent results as above starting with the decidability result of answering a query in ODRL0 (which does not include *not*[*policySet*]). The theorem in listing 9.1 states that for all environments, all single agreements, all subjects, all actions and all assets, either permission is granted, permission is denied, permission is unregulated or query is inconsistent.

Listing 9.1: Environments and Counts

```
Theorem queriesAreDecidable: ∀ (e:environment),
           ∀ (agr: agreement),
           ∀ (s:subject),
           ∀ (action:act),
           ∀ (a:asset),

(permissionGranted e [agr] s action a) \/
(permissionDenied e [agr] s action a)  \/
(queryInconsistent e [agr] s action a) \/
(permissionUnregulated e [agr] s action a).
```

We will then augment ODRL0 with the constructs we omitted from the full ODRL (resulting in what we have earlier called ODRL1 or ODRL2) including the troublesome construct *not[policySet]* and attempt to prove that the decidability results remain intact. There is a chance that a proof is not possible due to particulars of the Coq encoding we have used, in in which case, we will adjust our encoding.

9.3   SELinux

We started out by looking at DRELs and specifically ODRL where rights expressions are used to arbitrate access to assets under conditions. Recall that the main construct in ODRL is the agreement which specifies users, asset(s) and policies (as part of policy sets) whereby controls on users' access to the assets are described. This is reminiscent of how access control conditions are expressed in access control policy languages such as XACML and SELinux.

While XACML is a high-level and platform independent access control system SELinux is platform dependent (e.g. Linux based) and low-level. SELinux enhances the Discretionary access control (DAC) that most unix based systems employ by Mandatory access control (MAC) where designed access control policies are applied throughout the system possibly overriding whatever DAC is in place by the system users.

SELinux uses Linux's extended file attributes to attach a *security context* to passive entities (e.g. files, directories, sockets) and also to each active entity typically a Linux user space process. Security context is a data structure that is composed of a user, a role and a domain (or type). While users can map directly to ordinary user names they can also be defined separately. Roles are meant to group users and add flexibility when assigning permissions and are the basis for role-based access control (RBAC). Finally domains or types are the basis for defining common access control requirements for both passive and active entities.

The enforcement of SELinux policies are performed by the *security server*. Whenever a security operation is requested from user space by a system call, the security server is invoked to arbitrate the operation and either allow the operation or to deny it. Each

operation is identified by two pieces of information: an object class (e.g. file) and a permission (e.g. read, write). When an operation is requested to be performed on an object, the class and the permissions associated with the object along with security contexts of the source (typically the source entity is a process) and the object are passed to the security server. The security server consults the loaded policy (loaded at boot time) and allows or denies the access request [**?**].

9.4 SELinux Policy Language

The SELinux policy has four different kinds of statements: declarations, rules, constraints and assertions [archer]. Assertions are compile time checks that the *checkpolicy* tool performs at compile time. The other three kinds of statements however are evaluated at run-time.

Declaration statements declare a user, a role and a type.

Listing 9.2: Declarations

```
user u types Ru;

role r types Tr;

type t, attrib_{1}, ..., attrib_{n};
```

Rule statements define access vector rules and type transition rules. Access vector (AV) rules (see listing 9.3) specify which operations are allowed and whether to audit (log). Any operation not covered by AV rules are denied and all denied operations are logged. The semantics of the AV rule with avkind *allow* is: processes with type *sourcetype* are allowed to perform operations in *perm* on objects with class *obj-class* and type *targettype*. *auditallow* means to allow and audit, *dontaudit* means to never audit and finally *neverallow* provides a mechanism to override allow rules. When a process changes security context, the role may change, assuming a role transition rule exists relating the old and the new roles (listing **??**).

Listing 9.3: AV rule

```
avkind sourcetype targettype:object−class perm

avkind=allow, auditallow, dontaudit, neverallow
```

Constraints are additional conditions that must hold for an attempted operation to be allowed. Constrains relate all of their arguments (the security contexts) to the server (see listing 9.4. Whenever a permission is requested on an obj-class, the security server checks that the two security contexts are related by a constrain statement.

Listing 9.4: Constrain rule

```
constrain classes, perms, sourcetype, sourcerole, sourceuser, targettype, targetrole,
    targetuser
```

9.5 Agreements in SELinux

As with ODRL we will start by limiting the policy language to only allow AV rules. As mentioned earlier an operation not covered by a allow rule is denied by SELinux. We will make up explicit *deny* as required therefore an agreement is defined to be a combination of allow and deny rules. Allow and deny rules are mappings defined in listing 9.5.

Listing 9.5: allow/deny rule as a mapping

```
allow/deny rule :  T × (T × C) → 2^P
```

Listing 9.6: SELinux agreement

```
<agreement> ::= <avRule>
```

Listing 9.7: AV Rule

```
<avRule> ::=
   'allow' (T1, T2, C, P)                 ; allow rule
   'deny'  (T1, T2, C, P)                 ; deny rule
   'and'[ <avRule₁>, ..., <avRuleₘ> ]     ; conjunction
```

9.6 Environments

Environments are collections of *role-type* and *user-role* relations. A role-type relation $role(R, T)$ simply associates a role with a type. A user-role relation $user(U, R)$ associates a user with a role. An environment is consistent with respect to a security context $< T, R, U >$, if and only if $role(R, T)$ and $user(U, R)$ relations hold in the environment.

9.7 Queries in SELinux

The decision problem in SELinux access control is whether an entity with security context $< T1, R1, U1 >$ may perform action $P1$ to entity with object class $C1$ with security context $< T2, R2, U2 >$.

To answer such queries we use the authorization relation *auth(C, P, T1, R1, U1, T2, R2, U2)* which is equivalent to the *Permitted* predicate we saw earlier for ODRL.

$allow(T1, T2, C, P) \wedge (E\ consistent\ wrt\ < T1, R1, U1 >$ and $< T2, R2, U2 >)$
$\quad \implies auth(C, P, T1, R1, U1, T2, R2, U2)$

Listing 9.9: f_q^- for SELinux

$deny(T1, T2, C, P) \vee \neg(E\ consistent\ wrt\ < T1, R1, U1 >$ and $< T2, R2, U2 >)$
$\quad \implies \neg auth(C, P, T1, R1, U1, T2, R2, U2)$

9.8   Decidability of Queries in SELinux

In this thesis we will be investigating the question of decidability for answering queries in SELinux policies based on the same four outcomes we encountered earlier in **??** namely error, permitted, denied and "not applicable". We will first state a decidability theorem similar to theorem in listing 9.1 (minor adjustments may be needed to allow for differences with SELinux policy language) and present a proof for it in Coq. The literature in the SELinux implies only two outcomes are possible: permitted or denied. We will next attempt to prove this conjecture in Coq. Finally we will add constrain relations (see listing 9.10 and listing 9.11) to SELinux policies (which we have not included so far) and prove the same decidability results again for the augmented policy.

Listing 9.10: f_q^+ for SELinux

$allow(T1, T2, C, P) \wedge constrain(C, P, T1, R1, U1, T2, R2, U2) \wedge$
$(E\ consistent\ wrt\ < T1, R1, U1 >$ and $< T2, R2, U2 >)$
$\quad \implies auth(C, P, T1, R1, U1, T2, R2, U2)$

Listing 9.11: f_q^- for SELinux

$deny(T1, T2, C, P) \vee \neg\ constrain(C, P, T1, R1, U1, T2, R2, U2) \vee$
$\neg(E\ consistent\ wrt\ < T1, R1, U1 >$ and $< T2, R2, U2 >)$
$\quad \implies \neg auth(C, P, T1, R1, U1, T2, R2, U2)$

References

[1] Yves Bertot. Coq in a Hurry. Technical report, May 2010.

[2] Yves Bertot and Pierre Castéran. *Interactive Theorem Proving and Program Development. Coq'Art: The Calculus of Inductive Constructions.* Springer, 2004. http://www.labri.fr/perso/casteran/CoqArt/index.html.

[3] Robby Robson Geoff Collier, Harry Piccariello. A digital rights management ecosystem model for the education community. *DRM Whitepapers: Content Guard*, 2004.

[4] G. Huet and C. Paulin-mohring. The coq proof assistant reference manual. In *INRIA*, 2000.

[5] GÃľrard P. Huet. *The Gallina Specification language: A Case Study.* 1992.

[6] R. Iannella. Open digital rights language (odrl) version 1.1, 2002. [Online; accessed 09-February-2015].

[7] Robert C. Martin, Dirk Riehle, and Frank Buschmann. *Pattern Languages of Program Design 3.* Addison-Wesley, Reading, Mass., 1998.

[8] Riccardo Pucella and Vicky Weissman. A formal foundation for ODRL. *CoRR*, abs/cs/0601085, 2006.

[9] Beata Sarna-Starosta and Scott D. Stoller. Policy Analysis for Security-Enhanced Linux. In *Workshop on Issues in the Theory of Security*.

[10] Michael Carl Tschantz and Shriram Krishnamurthi. Towards reasonability properties for access-control policy languages. In *Proceedings of the Eleventh ACM Symposium on Access Control Models and Technologies*, SACMAT '06, pages 160–169, New York, NY, USA, 2006. ACM.

[11] Xin Wang, Guillermo Lao, Thomas DeMartini, Hari Reddy, Mai Nguyen, and Edgar Valenzuela. Xrml – extensible rights markup language. In *Proceedings of the 2002 ACM Workshop on XML Security*, XMLSEC '02, pages 71–79, New York, NY, USA, 2002. ACM.